

Data Mining Complete Exam Notes

CT 72502 (Elective I) · BEI / BEX IV-I · TU, IOE · tiers over 23 papers, 2069–2082 BS

3. Classification

12 hrs · in 23/23 papers · 7 topics, 3 TOP tier · largest chapter in the syllabus · numericals → Ch 3 companion.

Weight of the chapter:

- Decision Tree Classifier **TOP 9/23** → nearly always a full numerical.
- Bayesian Classifier **TOP 8/23**
- Issues: Overfitting, Validation, Model Comparison **TOP 7/23** → **confusion-matrix numerical appears in 8 papers.**
- ANN Classifier **HOT 5/23** · Rule Based **HOT 4/23**
- Basics **PIN 3/23** · Nearest Neighbor **PIN 2/23**

3.1 Basics and Algorithms

Why is classification a supervised learning method? **2 10**

PIN 4/23 82 Bh, 80 Bh, 74 Ch, 71 Shr

Classification

- Predicts **categorical** class labels (discrete, nominal).
- **Supervised**: the training tuples come w/ a known **class label**. The model learns from labelled answers.
- vs **unsupervised** (clustering): no labels given, groups discovered. → Ch 5
- vs **prediction / regression**: predicts a *continuous* value, not a label.

2-step process (**this is the block diagram Q**)

1. **Learning / model construction** (training)
 - Input: **training set**, each tuple w/ its class label.
 - Output: the model, as rules / tree / formula.
2. **Classification / model usage** (testing then deployment)
 - Estimate accuracy on an **independent test set**.
 - **Test set must be separate from the training set**, else the accuracy is optimistic (the model has memorised).
 - If accuracy acceptable → classify new, unlabelled data.

Where each classifier fits

- **Eager learners** build the model *before* seeing a query: decision tree, Bayes, ANN, rule-based, SVM.
- **Lazy learners** store the data and act only at query time: KNN, case-based.


Applications

- Credit / loan approval, fraud detection, medical diagnosis, target marketing, churn prediction, spam filtering.

Answer skeleton **10**

1. Definition + “supervised because labels are given”.
2. **Draw the 2-step block diagram** (training → model → testing → deployment).
3. Name the classifiers, 1 line each.
4. Close w/ the 2–3 metrics used to judge one.

The 2-step block diagram, drawn



(a) **Learning:** labelled training data → classification algorithm → model (here as IF-THEN rules). (b) **Classification:** the rules are checked on a separate test set, then applied to a new tuple (John Henry → risky). H&K Fig 8.1.

3.2 Decision Tree Classifier

What is a decision tree? How is information gain used for attribute selection? 8 10 4+6

TOP 13/23 82 Bh, 81 Bh, 79 Bh, 78 Bh, 76 Ch, 76 Ash, 75 Ch, 74 Ch, 73 Ch, 73 Shr, 72 Ka, 71 Ch, 70 Asa **Almost always a full tree-building numerical** → companion.

Structure

- Flow-chart-like **inverted tree**.
- **Root node:** topmost, the first test.
- **Internal node:** a test on one attribute.
- **Branch:** an outcome of that test.
- **Leaf node:** holds a **class label**.
- Path root → leaf = one classification rule.

Why popular

- No domain knowledge or parameter setting needed.
- Handles high-dimensional data.
- **Human-readable** → easy to convert to IF-THEN rules.
- Fast; good accuracy in most cases.

8.2 Decision Tree Induction 331




Figure 8.2 A decision tree for the concept *buys.computer*, indicating whether an *AllElectronics* customer is likely to purchase a computer. Each internal (nonleaf) node represents a test on

Hunt's algorithm 72 Ka 10 → the ancestor of all of them
 Let D_t = training records reaching node t , $y = \{y_1 \dots y_c\}$ the class labels.

1. If all records in D_t belong to the same class $y_t \rightarrow t$ is a **leaf** labelled y_t .
 2. If D_t is empty $\rightarrow$ leaf labelled w/ the **majority class of the parent**.
 3. If D_t has records of more than 1 class $\rightarrow$ pick an **attribute test**, split D_t into subsets, **recurse** on each subset.
- Descendants: **ID3** $\rightarrow$ **C4.5** $\rightarrow$ C5.0 (Quinlan), **CART**, SLIQ, SPRINT.

Generic induction algorithm

- **Greedy, top-down, recursive, divide-and-conquer.**
- Greedy = picks the locally best attribute at each node, never backtracks.
- **Stopping conditions:**
 - all tuples at the node share 1 class, or
 - no attributes left $\rightarrow$ leaf = majority class, or
 - no tuples left $\rightarrow$ leaf = majority class of parent.

The 3 attribute selection measures (memorise all three)

1. Information gain ID3 maximise

$$Info(D) = - \sum_{i=1}^m p_i \log_2(p_i)$$

$$Info_A(D) = \sum_{j=1}^v \frac{|D_j|}{|D|} \times Info(D_j)$$

$$Gain(A) = Info(D) - Info_A(D)$$

- $p_i = |C_{i,D}|/|D|$ = fraction of class i . $Info(D)$ = **entropy**, in bits.
- Pick the attribute w/ the **highest Gain**.
- **Bias:** prefers attrs w/ **many distinct values** (a unique ID gives $Info_A = 0$, so maximal gain, yet is useless).

2. Gain ratio C4.5 maximise

$$SplitInfo_A(D) = - \sum_{j=1}^v \frac{|D_j|}{|D|} \log_2\left(\frac{|D_j|}{|D|}\right)$$

$$GainRatio(A) = \frac{Gain(A)}{SplitInfo_A(D)}$$

- Normalises the gain $\rightarrow$ **fixes the many-values bias**.
- Unstable as $SplitInfo \rightarrow 0$; constrain $Gain$ to be at least the average.

Test conditions per attribute type 72 Ka asks this

- **Binary:** 2 outcomes, forced.
- **Nominal**
 - **Multi-way:** 1 branch per value.
 - **Binary:** split values into 2 subsets. $2^{v-1} - 1$ ways.
- **Ordinal:** same as nominal, but a grouping must **not break the order**.
- **Continuous**
 - **Binary:** test $A \leq t$ vs $A > t$.
 - **Multi-way:** discretise into ranges first.
 - **Split point:** sort values, try the **midpoint** of each adjacent pair, $\frac{a_i + a_{i+1}}{2}$; v values $\rightarrow v - 1$ candidate splits. Pick the one w/ minimum $Info_A(D)$.

3. Gini index CART minimise

$$Gini(D) = 1 - \sum_{i=1}^m p_i^2$$

$$Gini_A(D) = \frac{|D_1|}{|D|} Gini(D_1) + \frac{|D_2|}{|D|} Gini(D_2)$$

$$\Delta Gini(A) = Gini(D) - Gini_A(D)$$

- **Gini considers only BINARY splits** $\rightarrow$ for v values, try all $2^{v-1} - 1$ subset splits and keep the best.
- Pick the attribute w/ the **largest $\Delta Gini$** (equivalently smallest $Gini_A$).
- **Bias:** toward attrs w/ many values, and toward equal-sized balanced partitions.

Reading the measures

- All 3 measure **impurity**. Pure node $\rightarrow Info = 0$, $Gini = 0$.
- Max impurity: 2 classes 50/50 $\rightarrow Info = 1$, $Gini = 0.5$.

Overfitting + pruning 75 Ch 4

- Tree grows branches that fit **noise and outliers** → poor accuracy on new data.
- **Pre-pruning** (early stopping)
 - Halt splitting at a node; make it a leaf w/ the majority class.
 - Stop if: gain below a threshold, too few tuples, purity high enough.
 - **Hard to pick the threshold** → too high underfits, too low overfits.
- **Post-pruning** (grow full, then cut) → **preferred**
 - Remove subtrees from a fully grown tree, replace w/ a leaf.
 - CART uses *cost complexity* pruning w/ a separate prune set.
 - C4.5 uses *pessimistic* pruning, no extra data needed.
- Other fixes: more training data, cross-validation, ensembles.

Advantages

- Simple, fast to build and to apply.
- Self-explanatory, converts directly to rules.
- No normalisation or scaling needed.
- Handles both nominal and numeric attrs, and irrelevant attrs.

Disadvantages

- **Greedy** → locally optimal, not guaranteed globally optimal.
- Prone to **overfitting** without pruning.
- Unstable: a small data change can restructure the whole tree.
- Axis-parallel splits only → struggles w/ diagonal boundaries.
- Biased toward attrs w/ many values (unless gain ratio used).

3.3 Rule Based Classifier**How does a Rule Based Classifier work? Explain with example.** 7 2+8 9

HOT 5/23 80 Bh, 76 Ch, 74 Ash, 70 Ch, 69 Ch

Form

- Model = a set of **IF-THEN** rules.
- IF (age = youth) AND (student = yes) THEN buys_computer = 1
- **Antecedent** (LHS, the condition) → **Consequent** (RHS, the class).
- A rule **covers** a tuple if the antecedent holds; the rule is then said to be *triggered*.

Rule quality measures 70 Ch asks Accuracy + Laplace

verified/added

- n_{covers} = tuples covered by R , $n_{correct}$ = those it labels correctly, $|D|$ = size of the data set, k = no. of classes.

$$Coverage(R) = \frac{n_{covers}}{|D|} \quad Accuracy(R) = \frac{n_{correct}}{n_{covers}}$$

$$Laplace(R) = \frac{n_{correct} + 1}{n_{covers} + k} \quad m\text{-estimate}(R) = \frac{n_{correct} + kp}{n_{covers} + k}$$

- **Why Laplace**: plain accuracy rates a rule covering 1 tuple correctly as 1.0, which is meaningless. Laplace pulls small-coverage rules toward the prior, so a rule must cover *enough* tuples to score well.
- Eg 1 tuple, 1 correct, 2 classes: $Accuracy = 1/1 = 1.0$ but $Laplace = 2/3 = 0.67$.

Conflict resolution 76 Ch 2+4


- **Why needed**: a tuple may trigger **several rules predicting different classes, or no rule at all**.
- **Size ordering** (rule antecedent size): the *most specific* rule wins, ie the one w/ the most attribute tests. Intuition: more conditions = more evidence.
- **Class-based ordering**: classes sorted by decreasing prevalence (or by misclassification cost). All rules of one class come as a block.
- **Rule-based ordering (decision list)**: rules ranked by quality (accuracy, coverage, size) by the expert or the algorithm. **First rule that fires wins**.
- **No rule fires** → use a **default rule**, predicting the majority class of the tuples not covered by any rule. Always placed last.

Two useful properties

- **Mutually exclusive**: no 2 rules fire on the same tuple → no conflict possible.
- **Exhaustive**: every possible tuple is covered → no default rule needed.
- Rules extracted from a decision tree are **both** by construction.

Rule extraction from a decision tree

- 1 rule per **root-to-leaf path**. Each split condition ANDed; leaf = consequent.
- Mutually exclusive + exhaustive by construction.
- Rules can then be **pruned** individually, which may break exclusivity → then class-based ordering is applied.



One model three ways: (a) IF-THEN rules, (b) the decision tree they come from, one rule per root-to-leaf path, (c) a neural network. H&K Fig 1.9.

Decision tree classifier vs rule-based classifier 8

69 Ch “How is decision tree classifier different than rule based classifier?” **Note they are relatives, not rivals:** every tree converts to a rule set, so say that first and then give the differences. **verified/added**

Sequential covering (direct rule induction) 70 Ch asks CN2

verified/added

- Learns rules **directly from the data**, one class at a time.
- Algorithms: **CN2**, **RIPPER**, **AQ**, **FOIL**.
- Steps:
 1. Pick a class. Learn 1 rule that covers many of its tuples and few others.
 2. **Remove** the tuples that rule covers.
 3. Repeat on what remains until a stopping criterion is met.
- **CN2** specifically: general-to-specific beam search; starts from the empty antecedent and repeatedly appends the test that most improves the rule, scored by **entropy** and checked by a **likelihood-ratio significance test**. Produces an ordered rule list (decision list) + a default rule.

	Decision tree	Rule-based
Model form	One connected tree ; internal node = attribute test, leaf = class.	An unordered set or an ordered list of IF-THEN rules.
How it is built	Divide and conquer : recursively split the whole data set on the best attribute (Hunt's algorithm, ID3, C4.5, CART).	Separate and conquer (sequential covering): learn one rule, remove the tuples it covers, repeat on the remainder (CN2, RIPPER).
Coverage	Partitions the space exactly once : rules are automatically mutually exclusive and exhaustive , so every tuple lands in exactly one leaf.	Rules may overlap or leave gaps → needs a conflict-resolution strategy (size / class / rule ordering) and a default rule .
Attribute reuse	A test near the root constrains <i>every</i> rule below it, so an attribute chosen early is forced on all branches.	Each rule is grown independently, so rules need share no condition at all .
Pruning	Prunes <i>subtrees</i> ; removing a node changes every path through it.	Prunes <i>individual rules</i> , and may drop conditions from one rule without touching the rest. More flexible, but can break mutual exclusivity .
Replicated subtree problem	Suffers from it: the same subtree is duplicated under several branches when the concept is disjunctive.	Immune : one rule expresses the disjunct once.
Readability	Whole model visible in one diagram; easy to explain end to end, but a deep tree becomes unreadable.	Individual rules read very well on their own; a large unordered set is hard to see as a whole.

- **The link worth stating for the marks:** a decision tree can always be converted into a rule set by taking **one rule per root-to-leaf path**, ANDing the conditions along it. The reverse is *not* true, a general rule set need not correspond to any tree.
- So the real difference is **direct vs indirect induction**: rule-based classifiers learn rules straight from the data, decision trees learn a partition and get rules as a by-product.

3.4 Nearest Neighbor Classifier

What is a Nearest Neighbor Classifier? Main issues? Propose a fix. 1+3+4 3+6

PIN 4/23 81 Bh, 80 Ba, 80 Bh, 70 Asa 80 Bh is a numerical → companion.

Idea

- **Lazy learner** / instance-based: stores all training tuples, builds **no model**.
- All work happens at query time.
- “Tell me who your neighbours are and I will tell you who you are.”

Algorithm (KNN)

1. Choose k and a distance measure (usually Euclidean).
 2. Compute the distance from the query tuple to **every** training tuple.
 3. Take the k smallest → the k nearest neighbours.
 4. **Classification**: assign the **majority class** among those k .
Prediction: take their **mean value**.
- **Normalise the attributes first**, else large-range attrs dominate the distance.
 - k odd for 2 classes → avoids ties.

Issues (80 Ba asks exactly this)

- **Expensive at query time**: $O(n)$ distance computations per query, no model to shortcut it.
- **Curse of dimensionality**: distances lose meaning in high dims.
- **Scale sensitive**: needs normalisation.
- **Irrelevant attributes** distort distance equally w/ relevant ones.
- Needs the whole training set in memory.

How the distance metric is used in an instance-based classifier 10

70 Asa “What are the properties of a distance metric? How is distance metric used in instance based classifier?” **Two chapters in one question**: the properties are §2.4, the use is here. **The bridge between them is the answer.** verified/added


The four properties, restated §2.4

1. **Non-negativity**: $d(p, q) \geq 0$.
 2. **Identity**: $d(p, q) = 0$ iff $p = q$.
 3. **Symmetry**: $d(p, q) = d(q, p)$.
 4. **Triangle inequality**: $d(p, r) \leq d(p, q) + d(q, r)$.
- A function with all four is a **metric**. Euclidean, Manhattan and Minkowski qualify; **cosine similarity does not** (it is a similarity, not a distance).

- **Where the metric enters KNN, step by step**: compute $d(\text{query}, \text{every training tuple}) \rightarrow \text{sort} \rightarrow \text{take the } k \text{ smallest} \rightarrow$

Choosing k

- k too small → sensitive to **noise**, overfits.
- k too large → neighbourhood includes other classes, underfits.
- Pick by cross-validation; rule of thumb $k \approx \sqrt{n}$.



(a) 1-nearest neighbor (b) 2-nearest neighbor (c) 3-nearest neighbor

The query x with its (a) 1, (b) 2, (c) 3 nearest neighbours: the vote, and so the class, can change with k . Mishra's Ch 3 notes, p13.

Fixes / classifier to propose

- Speed: k-d trees, ball trees, partial distance, editing (drop redundant tuples).
- **Attribute weighting** or feature selection for irrelevant attrs.
- **Distance-weighted voting**: nearer neighbours count more, $w = 1/d^2$.
- **Propose an eager learner** instead: a **decision tree** or **Naive Bayes** builds the model once, then classifies in near-constant time, and ignores irrelevant attrs.

Why an instance-based classifier needs each one

- **The classifier is the distance function**. KNN stores the tuples and does no other modelling, so every decision it makes is a comparison of distances. A defective metric is a defective classifier.
- **Identity** → a training tuple identical to the query is at distance 0, so it is always among the k nearest. Without it a tuple could fail to match itself.
- **Symmetry** → “ q is near p ” and “ p is near q ” agree, so the neighbourhood is well defined however you sort.
- **Triangle inequality** → **this is the one that pays**: it lets **k-d trees** and **ball trees** discard an entire branch without measuring the tuples in it, turning an $O(n)$ scan into roughly $O(\log n)$. Drop it and every query must compare against the whole training set.
- **Non-negativity** → “nearest” means smallest, with a floor at 0.

majority vote (or distance-weighted vote, $w = 1/d^2$, which uses the actual magnitudes and not just the ordering).

- **Choice of metric changes the answer:** Euclidean for continuous attributes, **Manhattan** when attributes are independent scales, **Hamming** for nominal attributes, **Jaccard** for asymmetric binary. **Normalise first** — an attribute measured in thousands otherwise dominates the sum (§2.4).

3.5 Bayesian Classifier

Explain Naive Bayesian classification with a suitable example. 8 6 10 3+7

TOP 9/23 81 Ba, 79 Bh, 78 Bh, 76 Ash, 75 Ash, 72 Ch, 71 Ch, 71 Shr, 69 Ch **Numerical in 79 Bh and 72 Ch** → companion.

Bayes theorem

$$P(H|X) = \frac{P(X|H) P(H)}{P(X)}$$

- X = the tuple (evidence), H = hypothesis “ X belongs to class C ”.
- $P(H)$ = **prior**: probability of the class before seeing X .
- $P(H|X)$ = **posterior**: after seeing X . **This is what we want.**
- $P(X|H)$ = **likelihood**. $P(X)$ = **evidence**.
- 76 Ash asks prior vs posterior directly.

Naive Bayes

- Classify X to the class C_i w/ the largest $P(C_i|X)$.
- $P(X)$ is the **same for every class** → ignore it; just maximise $P(X|C_i)P(C_i)$.
- **“Naive” = class-conditional independence assumption:** given the class, the attributes are assumed independent, so

$$P(X|C_i) = \prod_{k=1}^n P(x_k|C_i)$$

- That turns 1 hard joint probability into n easy counts. **This is the whole trick.**

Algorithm steps (write these, then compute)

1. Compute the prior $P(C_i) = |C_{i,D}|/|D|$ for each class.
2. For each attribute value x_k in X , compute $P(x_k|C_i)$ from the counts.
 - Categorical: $P(x_k|C_i) = |\text{tuples in } C_i \text{ with } x_k|/|C_{i,D}|$.
 - Continuous: assume Gaussian, $g(x, \mu, \sigma) = \frac{1}{\sqrt{2\pi}\sigma} e^{-\frac{(x-\mu)^2}{2\sigma^2}}$.
3. Multiply: $P(X|C_i) = \prod_k P(x_k|C_i)$.
4. Compute $P(X|C_i) P(C_i)$ for every class.
5. **Predict the class w/ the maximum value.**

Laplacian correction 78 Bh short note 4

- **Problem:** if any $P(x_k|C_i) = 0$, the whole product collapses to **zero**, wiping out all the other evidence.
- **Fix:** add 1 to every count (and k to each denominator, k = no. of values):

$$P(x_k|C_i) = \frac{n_k + 1}{|C_{i,D}| + k}$$

- Eg counts 0, 990, 10 → probabilities 0, 0.990, 0.010; after correction $\frac{1}{1003}, \frac{991}{1003}, \frac{11}{1003} = 0.001, 0.988, 0.011$.
- **Estimates barely move, but the zero is gone.**

Bayesian Belief Network (BBN) 3+5 10

81 Ba, 71 Ch 81 Ba is a numerical → companion.


- **Limitation of Naive Bayes:** the independence assumption is **almost never true** in practice. Correlated attributes get double-counted, which skews the posterior.
- **BBN drops it:** allows a *subset* of attributes to be conditionally dependent.
- Two parts:
 1. **Directed acyclic graph (DAG):** node = variable, arc = direct causal dependence. Arc $Y \rightarrow Z$ means Y is a **parent** of Z .
 2. **Conditional probability table (CPT)** at each node, giving $P(Z | \text{parents}(Z))$.
- Joint probability: $P(x_1, \dots, x_n) = \prod_{i=1}^n P(x_i | \text{parents}(x_i))$
- A node is conditionally independent of its non-descendants given its parents.
- **When to use which** 71 Ch:
 - **Naive Bayes** → attributes roughly independent, need speed, little data, many attributes. Eg text/spam classification.
 - **BBN** → known causal structure and clear dependencies, need to reason w/ missing values. Eg medical diagnosis.

Advantages

- Fast to train and to classify; single pass over the data.
- Needs little training data.
- Handles missing values naturally (skip that factor).
- **Provably optimal** when the independence assumption actually holds.
- Incremental: counts can be updated as data arrives.

Disadvantages

- Independence assumption rarely holds → accuracy loss on correlated attrs.
- Zero-frequency problem (needs Laplacian correction).
- Cannot model interactions between attributes.
- Continuous attributes need a distribution assumption (usually Gaussian) or discretisation.



The 81 Ba network and its CPTs. Solved in the companion.

3.6 Artificial Neural Network Classifier

How does an ANN classifier work? Explain with algorithm. 8 2+6 5


HOT 6/23 82 Bh, 81 Ba, 80 Ba, 76 Ch, 75 Ch, 72 Ka 76 Ch is a gradient numerical → companion.

Structure

- Set of connected **units (neurons)** in layers, each connection carrying a **weight**.
- **Input layer** (1 unit per attribute) → one or more **hidden layers** → **output layer** (1 unit per class).
- **Multilayer feed-forward:** no cycles, no connections within a layer.
- Each unit: net input = $\sum_i w_i x_i + \theta$ (bias), then an **activation function**.
- **Sigmoid (logistic):**

$$\sigma(\theta) = \frac{1}{1 + e^{-\theta}}, \quad \sigma'(\theta) = \sigma(\theta)(1 - \sigma(\theta))$$

- **That derivative identity is why sigmoid is used** → makes backprop cheap.



The 76 Ch network. Gradients derived in the companion.

Backpropagation algorithm

1. **Initialise** weights and biases to small random numbers.
2. **Forward pass:** compute net input and output for each unit, layer by layer.

3. **Compute error at the output:** $Err_j = O_j(1 - O_j)(T_j - O_j)$ where $T_j = \text{target}$.
4. **Backpropagate:** for a hidden unit, $Err_j = O_j(1 - O_j) \sum_k Err_k w_{jk}$.
5. **Update weights and biases:** $w_{ij} += l \cdot Err_j \cdot O_i$, $\theta_j += l \cdot Err_j$ ($l = \text{learning rate}$).
6. **Repeat** until the terminating condition (error small, or epoch limit).

General considerations 72 Ka asks this

- **Network topology:** no. of hidden layers and units. Chosen by trial and error; no rule.
- **Input encoding:** normalise inputs to $[0, 1]$; one unit per nominal value.
- **Output encoding:** 1 unit for 2 classes, else 1 per class.
- **Learning rate l :** too large $\rightarrow$ oscillation; too small $\rightarrow$ slow. Often $1/t$ w/ epoch t .
- **Initial weights:** small random, to break symmetry.
- **Update mode:** case updating (per tuple) vs epoch updating (per pass).

Advantages

- High tolerance of **noisy** data.
- Classifies patterns it was not trained on.
- Works when the relationship is unknown / non-linear.
- Well suited to continuous-valued inputs and outputs.
- Inherently **parallel**.

Disadvantages

- **Black box:** poor interpretability, hard to justify a decision.
- Long training times.
- Many parameters set empirically (topology, l , epochs).
- Needs numeric input $\rightarrow$ encoding effort.
- Can get stuck in a **local minimum**.

3.7 Issues: Overfitting, Validation, Model Comparison

Confusion matrix: calculate accuracy, sensitivity, specificity, precision, recall. 4 10 3+5

TOP 11/23 82 Bh, 80 Ba, 79 Ba, 78 Bh, 75 Ash, 74 Ash, 74 Ch, 73 Ch, 73 Shr, 72 Ch, 70 Ch **8 papers. Highest-yield numerical in the whole syllabus** $\rightarrow$ companion.

The 4 building blocks

- **TP:** positive tuples correctly labelled positive.
- **TN:** negative tuples correctly labelled negative.
- **FP:** negative tuples wrongly labelled positive (*false alarm*).
- **FN:** positive tuples wrongly labelled negative (*miss*).
- $P = TP + FN$ (all actual positives), $N = FP + TN$ (all actual negatives).

		Predicted class		
		yes	no	
Actual class	yes	TP	FN	P
	no	FP	TN	N
Total		P'	N'	P + N

All the formulas (one table, learn it whole)

Measure	Formula
Accuracy	$(TP + TN) / (P + N)$
Error rate	$(FP + FN) / (P + N) = 1 - \text{accuracy}$
Sensitivity, recall, TPR	TP / P
Specificity, TNR	TN / N
False alarm rate, FPR	$FP / N = 1 - \text{specificity}$
Precision	$TP / (TP + FP)$
F_1 score	$\frac{2 \times \text{prec} \times \text{rec}}{\text{prec} + \text{rec}}$
F_β	$\frac{(1 + \beta^2) \text{prec} \times \text{rec}}{\beta^2 \text{prec} + \text{rec}}$

- **Recall = sensitivity = TPR.** Three names, one thing. Papers use all three.
- F_1 = harmonic mean of precision and recall.

THE trap in this chapter. Han & Kamber put **actual** on the rows, **predicted** on the columns. The IOE papers **flip this between years**: 74 Ch and 75 Ash write "Actual \ Predicted", but 78 Bh and 72 Ch write "Predicted / Actual".

If you read the orientation wrong, **FP and FN swap**, so sensitivity and specificity swap and precision is wrong. **Always copy the row/column labels onto your answer sheet first**, mark which axis is actual, then compute.

verified/added

When accuracy is the wrong measure 80 Ba 3

- **Class imbalance.** If 99% of tuples are negative, a classifier that always says “negative” scores **99% accuracy** and is useless.
- Examples: cancer screening, fraud detection, intrusion detection, defect detection.
- **Use instead:** sensitivity + specificity separately, precision/recall, F_1 , ROC/AUC, or cost-sensitive measures.
- Note: $accuracy = sensitivity \frac{P}{P+N} + specificity \frac{N}{P+N} \rightarrow$ accuracy hides the split.

Precision vs recall trade-off 74 Ch asks the inverse relation

- Loosen the threshold $\rightarrow$ more positives flagged $\rightarrow$ recall $\uparrow$, precision $\downarrow$.
- Tighten it $\rightarrow$ precision $\uparrow$, recall $\downarrow$.
- Perfect precision *and* recall is only possible w/ a perfect classifier $\rightarrow$ report F_1 to balance them.

ROC curve 73 Shr 1

- Plot of **TPR (y)** vs **FPR (x)** as the decision threshold sweeps.
- Diagonal = random guessing. Top-left corner = perfect.
- **AUC** = area under the curve. Closer to 1 = better; 0.5 = useless.
- Lets you compare 2 classifiers **independently of the threshold** and of class balance.

Evaluating a classifier: holdout, random subsampling, k-fold, bootstrap 8

79 Ba “Describe the working mechanism as well as the merits and demerits of the holdout method, random sampling, k-cross validation and bootstrap approaches in evaluating the performance of a classifier.” **8 marks means all four, each with mechanism + merit + demerit.** verified/added

Method	Working mechanism	Merits	Demerits
Holdout	Partition D once at random into a training set and a disjoint test set, typically $\frac{2}{3} / \frac{1}{3}$. Build the model on the first, measure accuracy on the second.	Simplest and cheapest : one model is built. Test set is genuinely unseen.	Estimate depends entirely on one lucky or unlucky split (high variance). Only part of the data trains, so accuracy is pessimistic on small D .
Random subsampling	Holdout repeated k times with a fresh random split each time; report the mean of the k accuracies.	Averages the split luck away $\rightarrow$ much more stable than one holdout. Simple to run.	k models to build. Splits overlap uncontrolled, so some tuples are tested many times and others never , and the runs are not independent.
k-fold cross-validation	Split D into k equal, mutually exclusive folds. In round i train on all folds but D_i and test on D_i . Accuracy = total correct over all k rounds / $ D $. $k = 10$ is the recommended value ; stratified folds keep each class's proportion.	Every tuple is tested exactly once and trains $k - 1$ times $\rightarrow$ uses the data fully and has low bias and low variance. The standard method, and the safe answer to “which would you use”.	Costs k model builds. Still slightly pessimistic (trains on $\frac{k-1}{k}$ of D). Leave-one-out ($k = n$) removes the bias but is prohibitively expensive and has high variance.
Bootstrap (.632)	Sample n tuples from D with replacement to form the training set; the tuples never drawn become the test set. A tuple is missed with probability $(1 - \frac{1}{n})^n \approx e^{-1} = 0.368$, so 63.2% of D trains on average. Repeat b times and combine: $acc = \frac{1}{b} \sum_{i=1}^b (0.632 acc_{test_i} + 0.368 acc_{train_i})$.	Best estimator for small data sets , where holdout wastes too much and folds are tiny. Reuses every tuple.	Optimistically biased , because duplicated training tuples also reappear in the model. The 0.632 weighting is a correction, not a cure. Most expensive of the four.

- **One-line ranking to close the answer:** holdout for speed, cross-validation for a trustworthy number, bootstrap when D is small.

Improving accuracy: ensembles (1 line each is enough)

- **Bagging:** train k models on bootstrap samples, take a **majority vote**. Reduces *variance*; helps unstable learners like decision trees.
- **Boosting / AdaBoost:** train models **sequentially**, each re-weighting the tuples the previous one got wrong; weighted

vote. Reduces *bias*; can overfit noise.

- **Random forest:** bagging + a random attribute subset at each split. Robust to overfitting, handles many attributes.

3.8 PYQ Mapping

N marks a question fully solved in the Ch 3 Numerical companion.

3.1 Basics

1. PIN 3/23 Block diagram of classification stages + inverse precision/recall relation + confusion matrix N. 2+3+5 74 Ch
2. Why is classification supervised? + impurity measures. 10 71 Shr
3. When do we use a classifier? (asked w/ the KNN numerical) 3 80 Bh

3.2 Decision Tree

4. N Draw a decision tree using ID3. 10 79 Bh
5. N Construct a tree using **information gain** + predict a new tuple. 8 78 Bh
6. N Which attribute would ID-3 choose as root? + draw the full tree. 2+5 76 Ch
7. TOP /23 Decision tree + how information gain selects attributes, w/ example. 8 73 Shr
8. Accuracy of classifiers + how to select the best root attribute. 4+6 76 Ash
9. PIN 2/23 Gini-Index w/ example + handling overfitting. 6+4 75 Ch 10 71 Ch
10. PIN 1/23 **Hunt's Algorithm** + test conditions per attribute type. 10 72 Ka
11. PIN 1/23 Decision tree w/ the concept of Naive Bayes classification. 10 74 Ch

3.3 Rule Based

12. HOT 4/23 How a rule-based classifier works, w/ example. 7 80 Bh 2+8 74 Ash
13. PIN 1/23 Rule-based classifier + **CN2** + define **Accuracy** and **Laplace** measures. 9 70 Ch
14. PIN 1/23 Why is conflict resolution necessary? + common strategies. 2+4 76 Ch

3.4 Nearest Neighbor

15. N Classify a flower w/ **KNN**, $K = 3$, Euclidean. 3+6 80 Bh
16. PIN 1/23 What is KNN? Main issues? Propose a classifier that solves them. 1+3+4 80 Ba

3.5 Bayesian

17. TOP /23 Naive Bayes w/ suitable example. 8 75 Ash 6 78 Bh
18. N Classify $X = (\text{Home Owner}=\text{No}, \text{Marital Status}=\text{Married}, \text{Income}=\$120K)$. 6 79 Bh
19. N Predict class for $X = (\text{age}=\text{youth}, \text{income}=\text{medium}, \text{student}=\text{yes}, \text{credit}=\text{fair})$. 10 72 Ch
20. PIN 1/23 Prior and posterior probabilities + algorithmic steps + strengths. 3+7 76 Ash
21. PIN 1/23 Bayes classifier + when to use Naive Bayes vs BBN. 10 71 Ch
22. PIN 1/23 Naive Bayes + solving overfitting in classification. 10 71 Shr
23. N Limitation of Naive Bayes + how BBN overcomes it + heart-disease inference. 3+5 81 Ba
24. PIN 1/23 Short note: **Laplacian correction**. 4 78 Bh

3.6 ANN

25. HOT 3/23 How does an ANN classifier work? w/ algorithm. 8 81 Ba, 75 Ch 2+6 72 Ka
26. N Compute activation at h_1 , $\partial E/\partial o$, $\partial E/\partial w_{12}$. 2+3+4 76 Ch
27. PIN 1/23 Short note: Neural Network Classifier. 5 80 Ba

3.7 Issues: Overfitting, Validation, Model Comparison

28. N TOP 7/23 **Confusion matrix numericals** (8 papers, same method, different tables):
 - Classification error, sensitivity, false alarm rate, specificity 3+5 80 Ba
 - Accuracy, TPR, FPR, precision 4 78 Bh
 - Accuracy, sensitivity, specificity, precision 10 75 Ash, 74 Ash
 - + recall 10 72 Ch
 - ROC + TPR, FPR, precision 1+3+6 73 Shr
 - Build the matrix from 2 predicted sequences, then 4 metrics 10 70 Ch
 - Given matrix $\rightarrow$ accuracy, sensitivity, precision 5 74 Ch

29. PIN 1/23 When can you not use accuracy? Give examples.

3 80 Ba

30. PIN 1/23 Short note: Overfitting problem in classification.

4 78 Bh

Papers added by the 2026-09 merge

The six papers of DM4.1BCT.pdf (2082 Bh, 2081 Bh, 2079 Ba, 2073 Ch, 2070 Asa, 2069 Ch), indexed here for the first time — the notes were built before that scan was merged. N = solved in this chapter's companion. verified/added

- N When do we use a decision tree? + a full ID-3 tree on the credit-risk table + how to evaluate the model. 2+9+3 82 Bh
- N Decision-tree algorithm + use ID3 to identify the root node only. 2+8 81 Bh
- Working mechanism and merits and demerits of holdout, random sampling, k -fold cross-validation and bootstrap. 8 79 Ba verified/added
- N Confusion matrix → accuracy, error rate, sensitivity, specificity, precision, recall. 4 79 Ba
- N What is ID3? + TPR, FPR and accuracy from a matrix that labels both axes Predicted. 2+6 73 Ch
- Importance of the homogeneity measure in a decision tree + explain Gini index. 8 70 Asa
- How is the distance metric used in an instance-based classifier? 10 70 Asa verified/added
- How is a decision tree classifier different from a rule-based classifier? 8 69 Ch verified/added
- Bayes' theorem + how Naive Bayes simplifies the computational complexity. 12 69 Ch
- Nearest neighbour classifier and when it is used 81 Bh · short notes Neural Network Classifier 82 Bh, Overfitting and ROC 73 Ch. 2+4 5 4

Coverage check

- All Ch 3 PYQs covered. 12 are numerical → all solved in the companion.
- Gaps filled by research verified/added :
 - Laplace and m-estimate rule measures and the CN2 algorithm 70 Ch: absent from every source note; taken from Tan/Steinbach and Han & Kamber §8.4.
 - Conflict resolution strategies 76 Ch: only named in the notes, not explained.
 - The confusion-matrix orientation trap (papers flip actual/predicted between years).

Sources used for this chapter

- Han & Kamber 3rd ed, ch 8 (basic concepts, decision tree induction, attribute selection measures, Bayes, rule-based, model evaluation) and ch 9 (BBN, backpropagation, lazy learners). Primary authority; every formula above checked against it.
- DBK Lectures 5–11 (Thapathali): Hunt's algorithm, tree construction, Naive Bayes, KNN, ANN, evaluation.
- data mining_/Chapter 3 - Classification.pdf, hunts algorithm.pdf, ClassificationDecisionTree.ppt, Chapter8_2.ppt, Bayes_Net.docx.

Abbreviations used in these notes

Only these are ever shortened. Anything not on this list is written out in full.

Short	Means	Short	Means
DM	data mining	DW	data warehouse
DB	database	DBMS	database management system
KDD	knowledge discovery in databases	DSS	decision support system
OLTP	on-line transaction processing	OLAP	on-line analytical processing
ROLAP	relational OLAP	MOLAP	multidimensional OLAP
HOLAP	hybrid OLAP	ETL	extract, transform, load
BI	business intelligence	ER	entity–relationship
attr(s)	attribute(s)	dim(s)	dimension(s)
multidim	multidimensional	txn(s)	transaction(s)
info	information	ops	operations
repo	repository	freq	frequent / frequency
config	configuration	no.	number
PCA	principal component analysis	SVD	singular value decomposition
SMC	simple matching coefficient	KNN	<i>k</i> -nearest neighbour
ANN	artificial neural network	POS	point of sale
FK	foreign key	UI	user interface
GIGO	garbage in, garbage out	w/, w/o	with, without
rm	remove	vs	versus
Eg	example / for example	+Fig	draw a figure
→	leads to / therefore	+	and

Exam-paper month codes (used in the year tags): Ba = Baishakh · Jth = Jestha · Asa = Ashar · Shr = Shrawan · Bh = Bhadra · Ash = Ashwin · Ka = Kartik · Mng = Mangsir · Po = Poush · Ma = Magh · Ch = Chaitra.

Tier chips: **TOP /23** = topic asked in 7 or more of the 17 papers · **HOT /23** = 4–6 papers · **PIN /23** = 1–3 papers but worth ≥ 8 marks. **Mark chips** such as **2+5** are the marks that question actually carried. **verified/added** marks anything researched and added that was not in the source notes.